

Dhruv Ragunathan

Contact Information

- LinkedIn: <https://www.linkedin.com/in/dhruv-ragunathan-908993b1/> (<https://www.linkedin.com/in/dhruv-ragunathan-908993b1/>)
- Github: <https://github.com/dragunat2016> (<https://github.com/dragunat2016>)

Presentation Date: November 20, 2023

Table of Contents

- [Overview](#)
- [Business Objectives](#)
- [Data Overview](#)
- [Data Preparation](#)
- [Exploratory Data Analysis](#)
- [Modeling](#)
- [Final Model Evaluation](#)
- [Recommendations](#)
- [Future Projects](#)
- [Reproduction Steps](#)

Overview



The Behavioral Risk Factor Surveillance System (BRFSS) is the nation's premier system of health-related telephone surveys that collects state data about U.S. residents regarding their health-related risk behaviors, chronic health conditions, and use of preventive services.

Established in 1984 with 15 states, BRFSS now collects data in all 50 states as well as the District of Columbia and three U.S. territories. BRFSS completes more than 400,000 adult interviews each year, making it the largest continuously conducted health survey system in the world.

Researchers have seen the opportunity to apply machine learning algorithms to make predictions on the data, since it was a feature rich dataset with hundreds-of-thousands of records.

Business Objectives



We have been tasked by the CDC to create models from previous BRFSS data that predicts diabetes. The CDC wants to help the people it surveys and alert them if they are at risk for diabetes given their survey results. Long-term the CDC would like to publish an application to Americans allowing them to fill out a form with questions on their vitals like BMI and habits such as exercise. Upon completing the form, the CDC would send back a diabetic risk to the person.

The motivation behind this is that diabetes is one of the most prevalent and costly diseases in the USA. Currently, 38 million people have diabetes of which 9 million are undiagnosed. When considering the precursor, prediabetes, that number jumps to 98 million people.

Diabetic patients are more likely to visit the emergency department and require expensive treatments and medications for their life. Reducing diabetes across the country would greatly improve the quality of life of millions of Americans.

Accuracy and precision are our primary metrics of evaluation. Accuracy defines the number of correct predictions made by the model over the total number of predictions. Precision defines the number of True positive identified over the true positive plus the false positive rate.

Optimizing on these two metrics should reduce the amount of false positives we encounter. We want to avoid false positives because they could result in unnecessary outreach and wasting resources. We will still record and review other metrics such as F1 score, ROC-AUC, and recall to review in-case these metrics are even for some models.

We will also be incorporating the "run time" of the model in our evaluation. Run time is the amount of time it takes to train and test the model.

A final model evaluation will be made by some heuristic combination of the accuracy, precision, and time it takes model too run. Any gains in accuracy and precision need to justify the time it takes to train and use the model.

Data Overview

Source

The 2015 data is available on this link from the CDC's website. The table with all the responses and the key donoting the data terms are also available. The link to the survey questions is [here \(https://www.cdc.gov/brfss/questionnaires/pdf-ques/2015-brfss-questionnaire-12-29-14.pdf\)](https://www.cdc.gov/brfss/questionnaires/pdf-ques/2015-brfss-questionnaire-12-29-14.pdf).

The page on the CDC's website containing the data is [here \(https://www.cdc.gov/brfss/annual_data/annual_data.htm\)](https://www.cdc.gov/brfss/annual_data/annual_data.htm).

The data on the CDC's page is in an ASCII format and hard too decode with time constraints. We found a CSV version of that data on Kaggle. The download link for the CSV is specifically [here \(https://www.kaggle.com/datasets/cdc/behavioral-risk-factor-surveillance-system\)](https://www.kaggle.com/datasets/cdc/behavioral-risk-factor-surveillance-system).

Full Link: <https://www.kaggle.com/datasets/cdc/behavioral-risk-factor-surveillance-system>
(<https://www.kaggle.com/datasets/cdc/behavioral-risk-factor-surveillance-system>)

Limitations

This is survey data where the user responses were segmented into several categories.

So the following limitations apply:

- Survey respondants may not be comfortable revealing sensitive information over the phone even if the response is anonymous.
- Many respondants who answer "no" for diabetes may actually have diabetes, but were not diagnosed. Note: That there was a significant imbalance of diabetes/pre-diabetes versus those who stated that they do not have the condition.
- Many variables that are continuous in nature were treated as ordinal in the study such as income and age. These variables were treated as ordinal as part of the models.

Data Preparation

The steps for data preparation and cleaning were done in this [notebook \(notebooks/Data_Cleaning.ipynb\)](#) for the sake of simplifying the main notebook.

This is the short version of the data cleaning process. For more detail please click the link above.

High - Level Process

- Selected for columns related to diabetes
- Dropped columns with significant data missing
- Reviewed the data in the features.
 - Values within features that corresponded to information like 'N/A', 'Refused', 'Didn't Know' were dropped.
 - Values were transformed to be more ordinal
- Combined Diabetes and Prediabetes data
- Addressed class imbalance by making the diabetes/non-diabetes records 50-50

```
In [6]: 1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 import warnings
6 warnings.filterwarnings("ignore")
7 import pickle
```

```
In [7]: 1 from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV, cross_val_score
2 from sklearn.preprocessing import StandardScaler, OneHotEncoder, FunctionTransformer
3 from sklearn.impute import SimpleImputer
4 from sklearn.compose import ColumnTransformer
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.svm import SVC
7 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
8 from sklearn.svm import LinearSVC
9 from sklearn.tree import DecisionTreeClassifier
10 from sklearn.naive_bayes import GaussianNB
11 from sklearn.neighbors import KNeighborsClassifier
12 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, plot_confusion_matrix, recall_score
13 from sklearn.compose import ColumnTransformer
14 from sklearn.pipeline import Pipeline
15 from sklearn import metrics
16 from xgboost import XGBClassifier
17 from datetime import datetime as dt
18 random_state=42
```

```
In [8]: 1 diab_df = pd.read_csv('diabetes_binary_5050_DR_BRFSS2015.csv')
2
3 diab_df.head()
```

Out[8]:

	Diabetes_binary	HighBP	Asthma	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack	PhysActivity	...	MentHlth	Employed
0	0.0	0.0	0.0	0.0	1.0	20.0	0.0	0.0	0.0	1.0	...	1.0	1.0
1	0.0	0.0	1.0	1.0	1.0	32.0	1.0	0.0	0.0	0.0	...	0.0	2.0
2	0.0	1.0	0.0	0.0	1.0	50.0	1.0	0.0	0.0	1.0	...	30.0	1.0
3	0.0	1.0	0.0	1.0	1.0	27.0	0.0	0.0	1.0	1.0	...	12.0	2.0
4	0.0	1.0	0.0	1.0	1.0	14.0	1.0	0.0	0.0	0.0	...	0.0	2.0

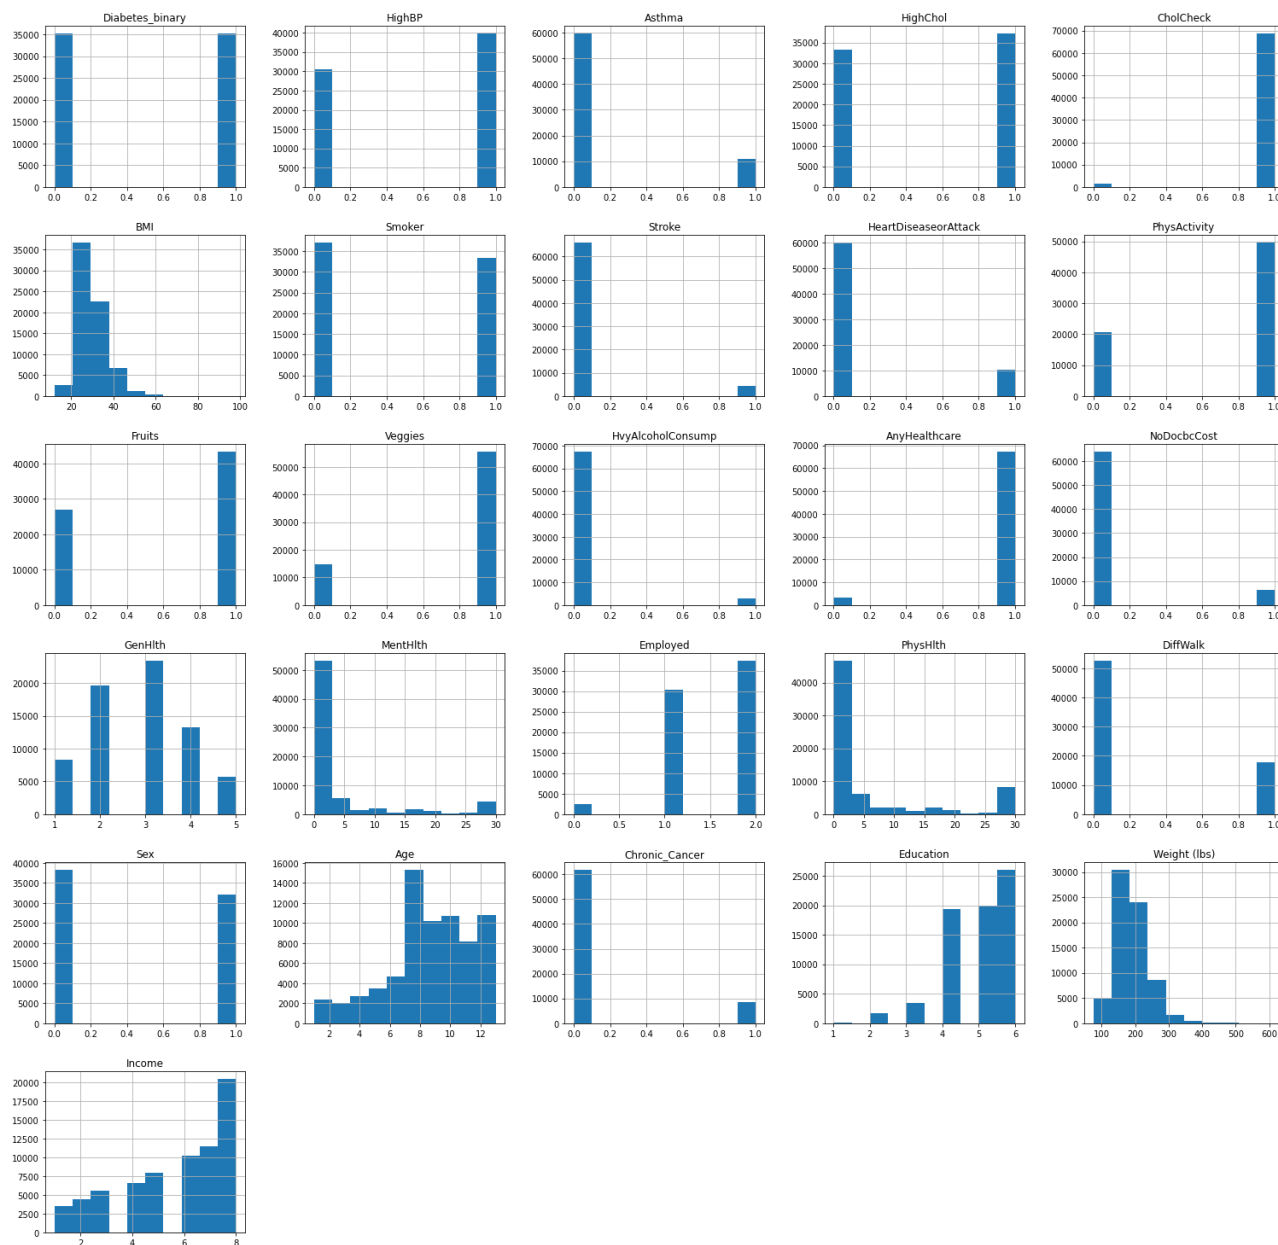
5 rows × 26 columns

In [9]: 1 diab_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 70252 entries, 0 to 70251
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Diabetes_binary                       70252 non-null  float64
1   HighBP                               70252 non-null  float64
2   Asthma                               70252 non-null  float64
3   HighChol                             70252 non-null  float64
4   CholCheck                            70252 non-null  float64
5   BMI                                  70252 non-null  float64
6   Smoker                               70252 non-null  float64
7   Stroke                               70252 non-null  float64
8   HeartDiseaseorAttack                 70252 non-null  float64
9   PhysActivity                         70252 non-null  float64
10  Fruits                               70252 non-null  float64
11  Veggies                              70252 non-null  float64
12  HvyAlcoholConsump                   70252 non-null  float64
13  AnyHealthcare                       70252 non-null  float64
14  NoDocbcCost                         70252 non-null  float64
15  GenHlth                             70252 non-null  float64
16  MentHlth                            70252 non-null  float64
17  Employed                            70252 non-null  float64
18  PhysHlth                            70252 non-null  float64
19  DiffWalk                            70252 non-null  float64
20  Sex                                  70252 non-null  float64
21  Age                                  70252 non-null  float64
22  Chronic_Cancer                      70252 non-null  float64
23  Education                           70252 non-null  float64
24  Weight (lbs)                        70252 non-null  float64
25  Income                              70252 non-null  float64
dtypes: float64(26)
memory usage: 13.9 MB
```

Exploratory Data Analysis

```
In [10]: 1 p = diab_df.hist(figsize = (26,26))
```



We can see a few interesting trends from the various histograms. First the diabetes versus non-diabetes is balanced as designed in the data cleaning process.

Second, High Blood pressure is also near balanced.


Weight is centered around near 200 points, which tracks on average.

There are more females than males in this study.

Higher incomes are mostly represented in the study. This could imply that the study is biased towards collecting data for those of a higher income. This would make sense since higher income individuals are more likely too have landlines.

Similarly, variables that show co-morbidities such as stroke, heart disease, and chronic cancer victims are not represented well in the data.

```
In [11]: 1 plt.figure(figsize=(50,50))
2 p = sns.heatmap(diab_df.corr(), annot=True,cmap = 'RdYlGn')
```



The vast majority of variables are not correlated with one another. This makes this data set could for modeling and less likely for overfitting/multicollinearity.

However, there is one exception. That being BMI and Weight. Since BMI is calculated from Weight this is not suprising.

To reduce the possibility of overfitting, we will drop the weight column. We chose to drop weight instead of BMI because BMI is more correlated with diabetes than weight is (0.29 vs 0.25). Therefore, dropping BMI as a feature would reduce the accuracy of the model more than weight would.

This data-driven decisions tracks with intuition. BMI is a better metric of determining how unhealthy an individual is since it incorporates height. A 6 foot individual weighing 180 pounds would be considered healthy while a 5 foot individual would not of that weight.

```
In [12]: 1 # Drop Weight column
          2
          3 diab_df = diab_df.drop('Weight (lbs)',axis=1)
```

```
In [13]: 1 diab_df.columns
```

```
Out[13]: Index(['Diabetes_binary', 'HighBP', 'Asthma', 'HighChol', 'CholCheck', 'BMI',
               'Smoker', 'Stroke', 'HeartDiseaseorAttack', 'PhysActivity', 'Fruits',
               'Veggies', 'HvyAlcoholConsump', 'AnyHealthcare', 'NoDocbcCost',
               'GenHlth', 'MentHlth', 'Employed', 'PhysHlth', 'DiffWalk', 'Sex', 'Age',
               'Chronic_Cancer', 'Education', 'Income'],
              dtype='object')
```

Modeling

Sections include

- Scaled Data for Model
- Ran Baseline Model
- Ran Additional Models
- Tuned best performing model from 'Additional Models' section
- Created a neural network since literature implied it was the best performing model for this use-case

Scaling Data

Using Standard Scaler to scale the data

```
In [14]: 1 sc_X = StandardScaler()
```

```
In [15]: 1 # Define Features and targets as X and y
          2
          3 X = diab_df.loc[:,diab_df.columns != 'Diabetes_binary']
          4 y = diab_df['Diabetes_binary']
```

```
In [16]: 1 X_scaled = sc_X.fit_transform(X)
          2 X_scaled = pd.DataFrame(X_scaled,columns=X.columns)
          3 X_scaled
```

```
Out[16]:
```

	HighBP	Asthma	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack	PhysActivity	Fruits	...	GenHlth
0	-1.14055	-0.425492	-1.057809	0.156285	-1.379308	-0.948568	-0.257453	-0.417718	0.647659	-1.265253	...	-0.751084
1	-1.14055	2.350221	0.945350	0.156285	0.301592	1.054221	-0.257453	-0.417718	-1.544023	0.790355	...	0.147490
2	0.87677	-0.425492	-1.057809	0.156285	2.822943	1.054221	-0.257453	-0.417718	0.647659	-1.265253	...	1.046064
3	0.87677	-0.425492	0.945350	0.156285	-0.398783	-0.948568	-0.257453	2.393962	0.647659	0.790355	...	0.147490
4	0.87677	-0.425492	0.945350	0.156285	-2.219758	1.054221	-0.257453	-0.417718	-1.544023	-1.265253	...	1.944638
...
70247	-1.14055	-0.425492	0.945350	0.156285	1.001967	-0.948568	-0.257453	-0.417718	-1.544023	-1.265253	...	1.046064
70248	-1.14055	-0.425492	0.945350	0.156285	-0.118633	1.054221	-0.257453	2.393962	-1.544023	0.790355	...	-0.751084
70249	0.87677	2.350221	0.945350	0.156285	-0.678933	-0.948568	-0.257453	2.393962	-1.544023	0.790355	...	1.944638
70250	0.87677	-0.425492	0.945350	0.156285	-1.659458	-0.948568	-0.257453	-0.417718	-1.544023	-1.265253	...	1.046064
70251	0.87677	-0.425492	0.945350	0.156285	-0.678933	-0.948568	-0.257453	2.393962	0.647659	0.790355	...	-0.751084

70252 rows × 24 columns

```
In [17]: 1 X_train, X_test, y_train, y_test = train_test_split(X_scaled,y, test_size=0.20)
```

```
In [76]: 1 # Pickle data to run models in other notebooks
2
3 with open('Variables/X_train.pickle', 'wb') as xtr:
4     pickle.dump(X_train,xtr)
5
6
```

```
In [77]: 1 #Store other variables
2
3
4
5 with open('Variables/X_test.pickle', 'wb') as xtst:
6     pickle.dump(X_test,xtst)
7
8 with open('Variables/y_train.pickle', 'wb') as ytr:
9     pickle.dump(y_train,ytr)
10
11
12 with open('Variables/y_test.pickle', 'wb') as ytst:
13     pickle.dump(y_test,ytst)
14
15
16
```

Baseline Model

- Start with baseline Logistic Regression
- Train data
- Make predictions from test data set
- Review metrics such as accuracy, recall, precision, ROC-AUC, and F1
- Review features

```
In [15]: 1 # Baseline Model is a logistic regression
2
3 lr_model = LogisticRegression()
4 lr_model.fit(X_train,y_train)
```

Out[15]: LogisticRegression()

```
In [16]: 1 lr_preds = lr_model.predict(X_train)
2
3 lr_train_acc = round(metrics.accuracy_score(y_train,lr_preds),3)
```

```
In [17]: 1 print('Training Accuracy score is ',lr_train_acc)
```

Training Accuracy score is 0.745


```
In [18]: 1 # Predictions from testing data set
2
3 y_pred = lr_model.predict(X_test)
```

```
In [19]: 1
2 lr_acc = metrics.accuracy_score(y_test, y_pred)
3 lr_rec = recall_score(y_test, y_pred)
4 lr_prec = precision_score(y_test, y_pred)
5 lr_roc_auc = roc_auc_score(y_test, y_pred)
6 lr_F1 = f1_score(y_test,y_pred)
7
8 print('Accuracy: ',lr_acc)
9 print('Recall: ',lr_rec)
10 print('Precision', lr_prec)
11 print('ROC - AUC',lr_roc_auc)
12 print('F1 Score',lr_F1)
```

Accuracy: 0.7494128531777098
Recall: 0.7691867943404316
Precision 0.7384742041712404
ROC - AUC 0.7494927450579391
F1 Score 0.7535176758837943

Let's take a look at the features this model prioritized.


```
In [20]: 1 fig = plt.figure(figsize = (10, 10))
2
3 feature_name = X_train.columns
4 coef_val = lr_model.coef_[0]
5
6 # creating the bar plot
7 plt.bar(feature_name, coef_val, color='blue',
8         width = 0.4)
9
10 plt.xlabel("Feature Name")
11 plt.ylabel("Coefficient Value")
12 plt.title("Magnitude of Logistical Regression Coefficients")
13 plt.xticks(rotation=90)
14 plt.show()
```



We can see that the feature given the most importance was GenHlth.


Other top features were High Blood Pressure, BMI, and Age.

Interestingly, heavy alcohol consumption did not positively affect diabetes correlation. Even though intuitively, one would think that more alcohol means more calories/sugar, which means higher likelihood for diabetes.

```
In [21]: 1 cm = confusion_matrix(y_test, y_pred)
2 lr_TN, lr_FP, lr_FN, lr_TP = confusion_matrix(y_test, y_pred).ravel()
3
4 print('True Positive(TP) = ', lr_TP)
5 print('False Positive(FP) = ', lr_FP)
6 print('True Negative(TN) = ', lr_TN)
7 print('False Negative(FN) = ', lr_FN)
```

```
True Positive(TP) = 5382
False Positive(FP) = 1906
True Negative(TN) = 5148
False Negative(FN) = 1615
```

```
In [22]: 1 cm_matrix = metrics.confusion_matrix(y_test,y_pred)
2
3 cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = cm_matrix, display_labels = [False, True])
4
5 cm_display.plot()
6 plt.show()
```



The confusion matrix above identifies similar amounts of true positives and true negatives. In addition, it also identified a similar number of false positives and false negatives.

- True Positive(TP) = 5320
- False Positive(FP) = 1982
- True Negative(TN) = 5033
- False Negative(FN) = 1716

These numbers are not too bad for a baseline model. The training and testing accuracy were similar, 74% indicating that the model is not overfitting the data. Let's see if we can use other models to improve these metrics from a baseline of 74%.

Additional Models

- Ran additional models such as Random Forest, XGB, Decision Tree Classifier, GaussianNB, and KNeighbors
- Reviewed metrics and selected one model for tuning
- Tuned the XGB model under hyper parameter tuning.
- Created confusion matrices, calculated metrics, and compared performance

```
In [23]: 1 # Models we want to test
2
3 model_arr = {}
4 model_arr['Logistical Regression'] = LogisticRegression()
5 model_arr['Random Forest'] = RandomForestClassifier()
6 model_arr['Decision Tree Classifier'] = DecisionTreeClassifier()
7 model_arr['XGB Classifier'] = XGBClassifier(gamma=0)
8 model_arr['SVC'] = SVC()
9 model_arr['GaussianNB'] = GaussianNB()
10 model_arr['KNeighbors'] = KNeighborsClassifier()
```

```
In [24]: 1 # loop over each classifier to evaluate performance
2
3 train_acc, test_acc, rec, prec, F1, Roc_Auc, trained_model, run_time = {}, {}, {}, {}, {}, {}, {}
4
5 for model_name in model_arr.keys():
6     model = model_arr[model_name]
7
8     start = dt.now()
9
10
11     # Fit the classifier
12     trained_model[model_name] = model.fit(X_train, y_train) #Store the trained model for further an
13
14     #Find training accuracy
15
16     y_train_pred = model.predict(X_train)
17
18     # Make predictions
19     y_pred = model.predict(X_test)
20
21     running_secs = (dt.now() - start).seconds
22
23     # Calculate metrics
24     train_acc[model_name] = accuracy_score(y_train, y_train_pred)
25     test_acc[model_name] = accuracy_score(y_test, y_pred)
26     rec[model_name] = recall_score(y_test, y_pred)
27     prec[model_name] = precision_score(y_test, y_pred)
28     F1[model_name] = f1_score(y_test, y_pred)
29     Roc_Auc[model_name] = roc_auc_score(y_test, y_pred)
30     run_time[model_name] = running_secs
```

```
In [25]: 1 measures = pd.DataFrame(index=model_arr.keys(), columns=['Training Accuracy', 'Testing Accuracy', 'Recall', 'Precision', 'F1 Score', 'Roc-AUC Score', 'Runtime (s)'])
2 measures['Training Accuracy'] = train_acc.values()
3 measures['Testing Accuracy'] = test_acc.values()
4 measures['Recall'] = rec.values()
5 measures['Precision'] = prec.values()
6 measures['F1 Score'] = F1.values()
7 measures['Roc-AUC Score'] = Roc_Auc.values()
8 measures['Runtime (s)'] = run_time.values()
9 measures
```

Out [25]:

	Training Accuracy	Testing Accuracy	Recall	Precision	F1 Score	Roc-AUC Score	Runtime (s)
Logistical Regression	0.745076	0.749413	0.769187	0.738474	0.753518	0.749493	0
Random Forest	0.997313	0.742652	0.780620	0.724115	0.751307	0.742805	7
Decision Tree Classifier	0.997331	0.655327	0.653566	0.654033	0.653799	0.655320	0
XGB Classifier	0.791214	0.748345	0.788767	0.728389	0.757376	0.748509	5
SVC	0.769719	0.753541	0.806917	0.727765	0.765300	0.753756	350
GaussianNB	0.714614	0.717743	0.703016	0.722638	0.712692	0.717683	0
KNeighbors	0.796107	0.709487	0.734886	0.697788	0.715857	0.709589	275

The disparity between the training and testing accuracy above for Random Forest and Decision Tree Classifier indicates that those models are highly overfit. Especially, the Decision Tree Classifier which had the lowest testing accuracy but a near 100% training accuracy.

The testing accuracies of the rest of the models were similar. SVC and XGB have the highest accuracies and have very close metrics to one another.

The only differences is that XGB has a marginally higher precision and SVC has a higher recall by 1% and a ROC-AUC score and accuracy score. Based on these metrics alone, it would make sense to chose SVC over XGB.

However, XGB runs significantly faster than SVC. In fact, XGB ran ~80 times faster than SVC. Note: Times may vary depending on machine. Since its significantly easier to use.

Ultimately, all of these models fall short of logistical's regressions accuracy to runtime ratio. Of all the models that ran in 0 seconds, logistical regression had the highest accuracy/precision.


That said, XGB has the potential to improve on these numbers through hyper-parameter tuning. We will be using this model for further analysis to try on improving on these results.

```
In [26]: 1 xgb = trained_model['XGB Classifier']
```

```
In [27]: 1 # Create a confusion matrix to visualize results
2
3 y_pred_xgb = xgb.predict(X_test)
4
5 cm = confusion_matrix(y_test, y_pred_xgb)
6 xgb_TN, xgb_FP, xgb_FN, xgb_TP = confusion_matrix(y_test, y_pred_xgb).ravel()
7
8 print('True Positive(TP) = ', xgb_TP)
9 print('False Positive(FP) = ', xgb_FP)
10 print('True Negative(TN) = ', xgb_TN)
11 print('False Negative(FN) = ', xgb_FN)
```

```
True Positive(TP) = 5519
False Positive(FP) = 2058
True Negative(TN) = 4996
False Negative(FN) = 1478
```

```
In [28]: 1 # Plot Results
2
3 xgb_cm_matrix = confusion_matrix(y_test, y_pred_xgb)
4
5 xgb_cm_display = ConfusionMatrixDisplay(confusion_matrix = xgb_cm_matrix, display_labels = [False, True])
6
7 xgb_cm_display.plot()
8 plt.show()
```



The confusion matrix above shows a high number of true positives/true negatives compared to the false positives/negatives. Let's see how many more correct prediction it made compared to the baseline model.

```
In [29]: 1 print('True Positive(TP) = ', xgb_TP)
2 print('False Positive(FP) = ', xgb_FP)
3 print('True Negative(TN) = ', xgb_TN)
4 print('False Negative(FN) = ', xgb_FN)
```

```
True Positive(TP) = 5519
False Positive(FP) = 2058
True Negative(TN) = 4996
False Negative(FN) = 1478
```


```
In [30]: 1 # Find the difference in correct predictions made between the xgboost Model and the Logistical Regres.
2 # Correct Predictions are defined as the number of TP + TN
3
4 lr_corr_pred = lr_TP + lr_TN # Correct number of predictions made by baseline logistic regression model
5 xgb_corr_pred = xgb_TP + xgb_TN # Correct number of predictions made by xgboost model
6
7 diff_preds_1 = abs(lr_corr_pred - xgb_corr_pred)
8
9 print("The xgboost model made", diff_preds_1, "more correct predictions than the baseline model.")
```

The xgboost model made 15 more correct predictions than the baseline model.

Let's take a look at what features XGB deemed important.

```
In [31]: 1 # Visualize Feature importance
2
3 pd.Series(xgb.feature_importances_, index=X_scaled.columns).sort_values(ascending=False).plot(kind='b
```

```
Out [31]: <AxesSubplot:>
```



Interestingly, the model put the highest weight on blood pressure by a significant margin. Almost 4 times higher than the next parameter of general health. This find tracks well with medical knowledge that high blood pressure and diabetes often are caused by unhealthy diet/health maintenance.

Hyper Parameter Tuning

Now that we have picked a model to further investigate, let's see if we can improve our accuracy through hyper parameter tuning. There are different methods for hyper parameter tuning such as grid searching and random search, but here we will use bayesian optimization. From research, we determined that this method is generally more successful than the others.

Due too time and resources constraints we will use this method instead of trying several and comparing the results.

Source: *Hyperparameter Optimization for Machine Learning Models Based on Bayesian Optimization*, Wu et. al. [link](https://www.sciencedirect.com/science/article/pii/S1674862X19300047) (<https://www.sciencedirect.com/science/article/pii/S1674862X19300047>)

```
In [32]: 1 from bayes_opt import BayesianOptimization
```

```
In [33]: 1 from sklearn.model_selection import cross_val_score
```

```
In [34]: 1 from hyperopt import fmin, tpe, hp
```

```
In [35]: 1 #Function that takes in parameters for xgboost and returns the highest roc-auc score in the cross val.
2
3 def xgboost_hyper_param(learning_rate,
4                           n_estimators,
5                           max_depth,
6                           subsample,
7                           gamma):
8
9     max_depth = int(max_depth)
10    n_estimators = int(n_estimators)
11
12    clf = XGBClassifier(
13        max_depth=max_depth,
14        learning_rate=learning_rate,
15        n_estimators=n_estimators,
16        gamma=gamma)
17    return np.mean(cross_val_score(clf, X_train, y_train, cv=3, scoring='roc_auc'))
18
19
```

```
In [36]: 1 # Parameters for xgboost model. Start with arbitrary parameter values
2
3 pbounds = {
4     'learning_rate': (0.01, 1.0),
5     'n_estimators': (100, 1000),
6     'max_depth': (3,10),
7     'subsample': (1.0, 1.0),
8     'gamma': (0, 5)}
9
```

```
In [37]: 1 #Instantiate the Optimizer
2
3 optimizer = BayesianOptimization(
4     f=xgboost_hyper_param,
5     pbounds=pbounds
6 )
```

```
In [38]: 1 optimizer.maximize(
2     init_points=2,
3     n_iter=3,
4 )
```

	iter	target	gamma	learni...	max_depth	n_esti...	subsample
	1	0.8072	1.012	0.4912	6.282	896.6	1.0
	2	0.8087	0.03199	0.2878	5.184	566.6	1.0
	3	0.8174	0.4783	0.353	5.156	567.6	1.0
	4	0.8272	2.406	0.3556	4.402	124.7	1.0
	5	0.7818	1.233	0.8737	6.646	715.8	1.0

```
In [39]: 1 optimizer.max
```

```
Out[39]: {'target': 0.8272277614471464,
'params': {'gamma': 2.4060556409414273,
'learning_rate': 0.3556184841702696,
'max_depth': 4.402240692590198,
'n_estimators': 124.67046120487794,
'subsample': 1.0}}
```

```
In [40]: 1 #parameters are in the 'params' keys
2
3 xgb_best_params = optimizer.max['params']
4
5 xgb_best_params
```

```
Out[40]: {'gamma': 2.4060556409414273,
'learning_rate': 0.3556184841702696,
'max_depth': 4.402240692590198,
'n_estimators': 124.67046120487794,
'subsample': 1.0}
```

```
In [41]: 1 # Create new classifier with the best params
2
3 gamma = xgb_best_params['gamma']
4 learning_rate = xgb_best_params['learning_rate']
5 max_depth = int(round(xgb_best_params['max_depth'])) # Needs to be an int not a float
6 n_estimators = int(round(xgb_best_params['n_estimators'])) # Needs to be an int not a float
7
8
9 xgb_tuned = XGBClassifier(gamma = gamma, learning_rate=learning_rate, max_depth=max_depth, n_estimators=
```

```
In [42]: 1 # Fit tuned model on training data
2
3 start = dt.now()
4
5 xgb_tuned.fit(X_train,y_train)
```

```
Out[42]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=2.4060556409414273,
gpu_id=-1, importance_type='gain', interaction_constraints='',
learning_rate=0.3556184841702696, max_delta_step=0, max_depth=4,
min_child_weight=1, missing=nan, monotone_constraints='()',
n_estimators=125, n_jobs=0, num_parallel_tree=1, random_state=0,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1.0,
tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [43]: 1 # Make predictions. Do the same for training data to determine if there is overfitting
2
3 y_pred_tuned_train = xgb_tuned.predict(X_train)
4 y_pred_tuned = xgb_tuned.predict(X_test)
5
6 running_secs_xgb = (dt.now() - start).seconds
```

```
In [44]: 1 xgb_tnd_trn_acc = accuracy_score(y_train,y_pred_tuned_train)
2 xgb_tnd_tst_acc = accuracy_score(y_test, y_pred_tuned)
3 xgb_tnd_rec = recall_score(y_test, y_pred_tuned)
4 xgb_tnd_rec_prec = precision_score(y_test, y_pred_tuned)
5 xgb_tnd_rec_roc_auc = roc_auc_score(y_test, y_pred_tuned)
6 xgb_tnd_rec_F1 = f1_score(y_test,y_pred_tuned)
7
8 print('Training Accuracy: ',xgb_tnd_trn_acc)
9 print('Testing Accuracy: ',xgb_tnd_tst_acc)
10 print('Recall: ',xgb_tnd_rec)
11 print('Precision', xgb_tnd_rec_prec)
12 print('ROC - AUC',xgb_tnd_rec_roc_auc)
13 print('F1 Score',xgb_tnd_rec_F1)
```

```
Training Accuracy: 0.757566591341791
Testing Accuracy: 0.7544658743149953
Recall: 0.7981992282406746
Precision 0.7326511871966418
ROC - AUC 0.7546425684724779
F1 Score 0.7640218878248974
```

It appears the parameters did not change the results significantly. For better visualization let's use a confusion matrix.

```
In [45]: 1 cm_xgb_tnd = confusion_matrix(y_test, y_pred_tuned)
2 TN_xgb_tnd, FP_xgb_tnd, FN_xgb_tnd, TP_xgb_tnd = confusion_matrix(y_test, y_pred_tuned).ravel()
3
4 print('True Positive(TP) = ', TP_xgb_tnd)
5 print('False Positive(FP) = ', FP_xgb_tnd)
6 print('True Negative(TN) = ', TN_xgb_tnd)
7 print('False Negative(FN) = ', FN_xgb_tnd)
```

```
True Positive(TP) = 5585
False Positive(FP) = 2038
True Negative(TN) = 5016
False Negative(FN) = 1412
```

```
In [46]: 1 # Difference between first tuning iteration and baseline model
2
3 lr_TP + lr_TN - TP_xgb_tnd - TN_xgb_tnd
```

```
Out[46]: -71
```

This tuning actually reduced the number of correct predictions the model makes.

```
In [47]: 1 # Add these values to our model dictionary
2 # Since pandas does not allow you to add rows without removing the indices correspond to the models,
3 # we need to recreate the table again
4
5 model_name = 'XGB Tuned 1'
6 model_arr['XGB Tuned 1'] = xgb_tuned
7
8 train_acc[model_name] = xgb_tnd_trn_acc
9 test_acc[model_name] = xgb_tnd_tst_acc
10 rec[model_name] = xgb_tnd_rec
11 prec[model_name] = xgb_tnd_rec_prec
12 F1[model_name] = xgb_tnd_rec_F1
13 Roc_Auc[model_name] = xgb_tnd_rec_roc_auc
14 run_time[model_name] = running_secs_xgb
```

```
In [48]: 1 measures = pd.DataFrame(index=model_arr.keys(), columns=['Training Accuracy', 'Testing Accuracy', 'Recall', 'Precision', 'F1 Score', 'Roc-AUC Score', 'Runtime (s)'])
2 measures['Training Accuracy'] = train_acc.values()
3 measures['Testing Accuracy'] = test_acc.values()
4 measures['Recall'] = rec.values()
5 measures['Precision'] = prec.values()
6 measures['F1 Score'] = F1.values()
7 measures['Roc-AUC Score'] = Roc_Auc.values()
8 measures['Runtime (s)'] = run_time.values()
9 measures
```

Out[48]:

	Training Accuracy	Testing Accuracy	Recall	Precision	F1 Score	Roc-AUC Score	Runtime (s)
Logistical Regression	0.745076	0.749413	0.769187	0.738474	0.753518	0.749493	0
Random Forest	0.997313	0.742652	0.780620	0.724115	0.751307	0.742805	7
Decision Tree Classifier	0.997331	0.655327	0.653566	0.654033	0.653799	0.655320	0
XGB Classifier	0.791214	0.748345	0.788767	0.728389	0.757376	0.748509	5
SVC	0.769719	0.753541	0.806917	0.727765	0.765300	0.753756	350
GaussianNB	0.714614	0.717743	0.703016	0.722638	0.712692	0.717683	0
KNeighbors	0.796107	0.709487	0.734886	0.697788	0.715857	0.709589	275
XGB Tuned 1	0.757567	0.754466	0.798199	0.732651	0.764022	0.754643	5

Let's see if we can further improve them by increasing the bounds and also by increasing the number of iterations the optimizer runs over.

```
In [49]: 1 pbounds2 = {
2     'learning_rate': (0.01, 0.6),
3     'n_estimators': (100, 300),
4     'max_depth': (3, 7),
5     'subsample': (1.0, 1.0), # Won't allow values over 1.0
6     'gamma': (5, 20)}
7
8
9 optimizer2 = BayesianOptimization(
10     f=xgboost_hyper_param,
11     pbounds=pbounds2
12 )
```

```
In [50]: 1 # Increased init_points and n_iter by 1 from previous tuning
2
3 optimizer2.maximize(
4     init_points=4,
5     n_iter=5,
6 )
```

	iter	target	gamma	learn...	max_depth	n_esti...	subsample
	1	0.7889	1.224	0.8298	6.204	706.0	1.0
	2	0.8289	4.624	0.1121	5.794	705.2	1.0
	3	0.813	1.293	0.7179	5.142	988.3	1.0
	4	0.7824	0.9487	0.6586	9.097	517.2	1.0
	5	0.8232	3.325	0.3872	7.412	790.9	1.0
	6	0.8287	4.448	0.1623	3.624	635.5	1.0
	7	0.8267	3.234	0.5833	3.322	531.9	1.0
	8	0.8162	0.3406	0.8856	3.381	645.7	1.0
	9	0.796	1.448	0.4718	9.502	782.1	1.0


```
In [51]: 1 xgb_best_params2 = optimizer2.max['params']
2 xgb_best_params2
```

```
Out[51]: {'gamma': 4.624140496883825,
'learning_rate': 0.11214395671681449,
'max_depth': 5.794307663328713,
'n_estimators': 705.2226048065555,
'subsample': 1.0}
```

```
In [52]: 1 # Create new classifier with the best params
2 gamma = xgb_best_params2['gamma']
3 learning_rate = xgb_best_params2['learning_rate']
4 max_depth = int(round(xgb_best_params2['max_depth'])) # Needs to be an int not a float
5 n_estimators = int(round(xgb_best_params2['n_estimators'])) # Needs to be an int not a float
6
7
8 xgb_tuned_2 = XGBClassifier(gamma = gamma, learning_rate=learning_rate, max_depth=max_depth, n_estimators=
```

```
In [53]: 1 start = dt.now()
2
3 xgb_tuned_2.fit(X_train, y_train)
```

```
Out[53]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=4.624140496883825,
gpu_id=-1, importance_type='gain', interaction_constraints='',
learning_rate=0.11214395671681449, max_delta_step=0, max_depth=6,
min_child_weight=1, missing=nan, monotone_constraints=()),
n_estimators=705, n_jobs=0, num_parallel_tree=1, random_state=0,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1.0,
tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [54]: 1 y_pred_tuned_2_train = xgb_tuned_2.predict(X_train)
2 y_pred_tuned_2 = xgb_tuned_2.predict(X_test)
3
4 running_secs_xgb_2 = (dt.now() - start).seconds
```

```
In [55]: 1 xgb_tnd_2_trn_acc = accuracy_score(y_train, y_pred_tuned_2_train)
2 xgb_tnd_2_tst_acc = accuracy_score(y_test, y_pred_tuned_2)
3 xgb_tnd_2_rec = recall_score(y_test, y_pred_tuned_2)
4 xgb_tnd_2_rec_prec = precision_score(y_test, y_pred_tuned_2)
5 xgb_tnd_2_rec_roc_auc = roc_auc_score(y_test, y_pred_tuned_2)
6 xgb_tnd_2_rec_F1 = f1_score(y_test, y_pred_tuned_2)
7
8 print('Training Accuracy: ', xgb_tnd_2_trn_acc)
9 print('Testing Accuracy: ', xgb_tnd_2_tst_acc)
10 print('Recall: ', xgb_tnd_2_rec)
11 print('Precision', xgb_tnd_2_rec_prec)
12 print('ROC - AUC', xgb_tnd_2_rec_roc_auc)
13 print('F1 Score', xgb_tnd_2_rec_F1)
```

```
Training Accuracy: 0.7644347965338695
Testing Accuracy: 0.7538965198206533
Recall: 0.7974846362726883
Precision 0.7321873769846476
ROC - AUC 0.7540726271808579
F1 Score 0.7634423313722808
```

```
In [56]: 1 # Add these values to our model dictionary
2 # Since pandas does not allow you to add rows without removing the indices correspond to the models,
3 # we need to recreate the table again
4
5 model_name = 'XGB Tuned 2'
6 model_arr['XGB Tuned 2'] = xgb_tuned_2
7
8 train_acc[model_name] = xgb_tnd_2_trn_acc
9 test_acc[model_name] = xgb_tnd_2_tst_acc
10 rec[model_name] = xgb_tnd_2_rec
11 prec[model_name] = xgb_tnd_2_rec_prec
12 F1[model_name] = xgb_tnd_2_rec_F1
13 Roc_Auc[model_name] = xgb_tnd_2_rec_roc_auc
14 run_time[model_name] = running_secs_xgb_2
```

```
In [57]: 1 measures = pd.DataFrame(index=model_arr.keys(), columns=['Training Accuracy', 'Testing Accuracy', 'Recall', 'Precision', 'F1 Score', 'Roc-AUC Score', 'Runtime (s)'])
2 measures['Training Accuracy'] = train_acc.values()
3 measures['Testing Accuracy'] = test_acc.values()
4 measures['Recall'] = rec.values()
5 measures['Precision'] = prec.values()
6 measures['F1 Score'] = F1.values()
7 measures['Roc-AUC Score'] = Roc_Auc.values()
8 measures['Runtime (s)'] = run_time.values()
9 measures
```

Out[57]:

	Training Accuracy	Testing Accuracy	Recall	Precision	F1 Score	Roc-AUC Score	Runtime (s)
Logistical Regression	0.745076	0.749413	0.769187	0.738474	0.753518	0.749493	0
Random Forest	0.997313	0.742652	0.780620	0.724115	0.751307	0.742805	7
Decision Tree Classifier	0.997331	0.655327	0.653566	0.654033	0.653799	0.655320	0
XGB Classifier	0.791214	0.748345	0.788767	0.728389	0.757376	0.748509	5
SVC	0.769719	0.753541	0.806917	0.727765	0.765300	0.753756	350
GaussianNB	0.714614	0.717743	0.703016	0.722638	0.712692	0.717683	0
KNeighbors	0.796107	0.709487	0.734886	0.697788	0.715857	0.709589	275
XGB Tuned 1	0.757567	0.754466	0.798199	0.732651	0.764022	0.754643	5
XGB Tuned 2	0.764435	0.753897	0.797485	0.732187	0.763442	0.754073	38

These numbers seem slightly better than the initial Xgboost model as well as the baseline. Given the magnitude of data we are working over, over 10,000 records, the gains are marginal at best.


However, interestingly it took less time for the second tuned model to evaluate. The second tuning showed good improvements.

With more computation resources, it would be interesting to see how much higher we can increase the accuracy of the model.

```
In [58]: 1 cm_xgb_tnd = confusion_matrix(y_test, y_pred_tuned_2)
2 TN_xgb_tnd2, FP_xgb_tnd2, FN_xgb_tnd2, TP_xgb_tnd2 = confusion_matrix(y_test, y_pred_tuned_2).ravel()
3
4 print('True Positive(TP) = ', TP_xgb_tnd2)
5 print('False Positive(FP) = ', FP_xgb_tnd2)
6 print('True Negative(TN) = ', TN_xgb_tnd2)
7 print('False Negative(FN) = ', FN_xgb_tnd2)
```

```
True Positive(TP) = 5580
False Positive(FP) = 2041
True Negative(TN) = 5013
False Negative(FN) = 1417
```

```
In [59]: 1 cm_xgb_tnd = confusion_matrix(y_test, y_pred_tuned_2)
2
3 cm_xgb_tnd = ConfusionMatrixDisplay(confusion_matrix = cm_xgb_tnd, display_labels = [False, True])
4
5 cm_xgb_tnd.plot()
6 plt.show()
```



The tuned model has performed slightly better than the baseline and initial XGBoost model. Since the percentages are small, let's see how many correct predictions this translates too.

```

In [60]: 1 # Find the difference in correct predictions made between the tuned XGBoost Model and the un-tuned model
2 # Correct Predictions are defined as the number of TP + TN
3
4 lr_corr_pred = lr_TP + lr_TN # Correct number of predictions made by baseline logistic regression model
5 xgb_corr_pred = xgb_TP + xgb_TN # Correct number of predictions made by XGBoost model
6
7 xgb_tnd_corr_pred = TP_xgb_tnd + TN_xgb_tnd
8
9 diff_preds_1 = xgb_corr_pred - lr_corr_pred
10 diff_preds_2 = xgb_tnd_corr_pred - xgb_corr_pred
11 diff_preds_3 = xgb_tnd_corr_pred - lr_corr_pred
12
13
14 print("The initial XGBoost model made",diff_preds_1,"more correct predictions than the baseline model")
15 print("The tuned XGBoost model made",diff_preds_2,"more correct predictions than the initial XGBoost model")
16 print("The tuned XGBoost model made",diff_preds_3,"more correct predictions than the baseline model.")
17
18

```

The initial XGBoost model made -15 more correct predictions than the baseline model.
The tuned XGBoost model made 86 more correct predictions than the initial XGBoost model.
The tuned XGBoost model made 71 more correct predictions than the baseline model.

Through our iterative modeling process we are increasing the accuracy of our model. However, these increases are marginal at best over a dataset that has tens of thousands of values.


It's unclear if the time and effort spent on tuning the model is worth the gain in accuracy.

```

In [61]: 1 pd.Series(xgb_tuned_2.feature_importances_, index=X_scaled.columns).sort_values(ascending=False).plot


```

Out[61]: <AxesSubplot:>



There does not seem to be a huge difference in the features. Though the coefficient for HighBP increased and the rest decreased.

```
In [62]: 1 x_axis = ["Logistical Regression", "XGBoost", "XGBoost with Parameter Tuning"]
2 y_axis = [lr_corr_pred, xgb_corr_pred, xgb_tnd_corr_pred]
3
4 plt.plot(x_axis,y_axis,linewidth = 3)
5 plt.xlabel('Model')
6 plt.ylabel('Correct Predictions')
7 plt.title('Number of Correct Predictions over Models')
8
9 # Show the plot
10 plt.show()
```



```
1 ## Using Neural Network Models on the Data
2
3 In addition to our own iterative modeling, we wanted to research the techniques experts were finding
4 to be the most accurate in predicting diabetes.
5
6 We found several articles that found neural networks to provide the best model including one that
7 used a dataset from a previous BRFSS dataset in a previous year.
8
9 The following sources evaluated the implementing different machine learning models on diabetes data.
10 They concluded that neural networks were the best model when evaluating based on accuracy.
11
12 * Building Risk Prediction Models for Type 2 Diabetes Using Machine Learning Techniques*, Xie et.
13 al. <a href="https://www.cdc.gov/pcd/issues/2019/19_0109.htm">link</a>
14 * This article used the 2014 data from the survey to create these models.
15 * Cardiovascular complications in a diabetes prediction model using machine learning: a systematic
16 review*, Kee et. al. <a href="https://link.springer.com/article/10.1186/s12933-023-01741-7">link</a>
17
18 We created our own neural network based on the data. Due to the size and amount of text generated by
19 neural networks, we ran them on a different notebook. We saved the best model and loaded it here to
20 create the confusion matrix, graphs, etc.
21
22 The analysis and notebook containing the optimization of the neural network is <a
23 href="notebooks/Neural_Network_Modeling.ipynb">here</a>
24
25 Only the architecture for the final model was included in the notebook.
26
27 The neural networks architecture is:
28
29 Neural
30 * 3 dense layers
31 * 40 neurons in the first layer
32 * 20 neurons in the second
33 * 10 neurons in the third
34 * relu activation
35 * Use sigmoid curve
36 * Early Stopping
```

```
In [1]: 1 import keras
2 from keras import models
3 from keras.models import Sequential
4 from keras.layers import Dense
5 import tensorflow as tf
6 from keras import callbacks
7 from keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
In [64]: 1 # Uncomment if not running from scratch.
2
3 #nn_model = keras.models.load_model('Neural_Network')
```

```
In [18]: 1 # Instantiate the model
2
3 nn_model = Sequential()
4 num_features = X_train.shape[1]
```

```
In [19]: 1 # 1st layer: input_dim=8, 40 nodes, RELU
2 nn_model.add(Dense(40, input_dim=num_features, activation='relu'))
3 # 2nd layer: 20 nodes, RELU
4 nn_model.add(Dense(20, activation='relu'))
5 # 3rd layer:
6 nn_model.add(Dense(10, activation='relu'))
7
8 # output layer: dim=1, activation sigmoid
9 nn_model.add(Dense(1, activation='sigmoid' ))
10
11 # early stopping - monitor for validation loss. Wait for 5 epochs if loss increases. Allow for 1% er
12 es = [EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=5,min_delta=1),
13       ModelCheckpoint(filepath='Neural_Network', monitor='val_loss',
14                       save_best_only=True)]
15
16 # Compile the model
17 nn_model.compile(loss='binary_crossentropy', # since we are predicting 0/1
18                 optimizer='adam',
19                 metrics=['accuracy'])
```

```
In [20]: 1 history = nn_model.fit(X_train,
2                             y_train,
3                             validation_data=(X_test, y_test),
4                             epochs=30,
5                             batch_size=16,
6                             callbacks=es)
```

Epoch 1/30

3490/3513 [=====>.] - ETA: 0s - loss: 0.5167 - accuracy: 0.7441WARNING:tensorflow:From /Users/dhruvruganathan/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/tensorflow/python/training/tracking/tracking.py:111: Model.state_updates (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.

Instructions for updating:

This property should not be used in TensorFlow 2.0, as updates are applied automatically.

WARNING:tensorflow:From /Users/dhruvruganathan/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/tensorflow/python/training/tracking/tracking.py:111: Layer.updates (from tensorflow.python.keras.engine.base_layer) is deprecated and will be removed in a future version.

Instructions for updating:

This property should not be used in TensorFlow 2.0, as updates are applied automatically.

INFO:tensorflow:Assets written to: Neural_Network/assets

3513/3513 [=====] - 6s 2ms/step - loss: 0.5164 - accuracy: 0.7443 - val_loss: 0.5084 - val_accuracy: 0.7498

Epoch 2/30

3480/3513 [=====>.] - ETA: 0s - loss: 0.5043 - accuracy: 0.7519INFO:tensorflow:Assets written to: Neural_Network/assets

3513/3513 [=====] - 5s 1ms/step - loss: 0.5047 - accuracy: 0.7515 - val_loss: 0.5084 - val_accuracy: 0.7481

Epoch 3/30

3486/3513 [=====>.] - ETA: 0s - loss: 0.5019 - accuracy: 0.7546INFO:tensorflow:Assets written to: Neural_Network/assets

3513/3513 [=====] - 5s 2ms/step - loss: 0.5022 - accuracy: 0.7545 - val_loss: 0.5070 - val_accuracy: 0.7528

Epoch 4/30

3477/3513 [=====>.] - ETA: 0s - loss: 0.5010 - accuracy: 0.7544INFO:tensorflow:Assets written to: Neural_Network/assets

3513/3513 [=====] - 5s 1ms/step - loss: 0.5009 - accuracy: 0.7546 - val_loss: 0.5053 - val_accuracy: 0.7519

Epoch 5/30

3498/3513 [=====>.] - ETA: 0s - loss: 0.4994 - accuracy: 0.7560INFO:tensorflow:Assets written to: Neural_Network/assets

3513/3513 [=====] - 5s 1ms/step - loss: 0.4993 - accuracy: 0.7561 - val_loss: 0.5042 - val_accuracy: 0.7517

Epoch 6/30

3513/3513 [=====] - 5s 1ms/step - loss: 0.4985 - accuracy: 0.7564 - val_loss: 0.5048 - val_accuracy: 0.7511

Epoch 00006: early stopping

```
In [21]: 1 # Predict on training data
2 # The data of y_preds_nn is float not binary 0/1 so we cannot compare it to y_test in current state
3
4 y_pred_nn = nn_model.predict(X_test)
5
6 y_pred_nn
```

```
Out[21]: array([[0.8572431 ],
                [0.9653896 ],
                [0.00631964],
                ...,
                [0.648824  ],
                [0.37289095],
                [0.7267164 ]], dtype=float32)
```

```
In [22]: 1 # We will round y_preds_nn to 0 or 1 depending on if it's above or below 0.5
2
3 y_pred_nn_rnd = np.around(y_pred_nn,0)
4
5 y_pred_nn_rnd
```

```
Out[22]: array([[1.],
                [1.],
                [0.],
                ...,
                [1.],
                [0.],
                [1.]], dtype=float32)
```

```
In [23]: 1 # Calculate metrics below
2
3 nn_trn_acc = 0.7578 # Pulled from neural network notebook
4 nn_tst_acc = accuracy_score(y_test, y_pred_nn_rnd)
5 nn_rec = recall_score(y_test, y_pred_nn_rnd)
6 nn_rec_prec = precision_score(y_test, y_pred_nn_rnd)
7 nn_rec_roc_auc = roc_auc_score(y_test, y_pred_nn_rnd)
8 nn_rec_F1 = f1_score(y_test, y_pred_nn_rnd)
9
10 print('Training Accuracy: ',nn_trn_acc)
11 print('Testing Accuracy: ',nn_tst_acc)
12 print('Recall: ',nn_rec)
13 print('Precision', nn_rec_prec)
14 print('ROC - AUC',nn_rec_roc_auc)
15 print('F1 Score',nn_rec_F1)
```

```
Training Accuracy: 0.7578
Testing Accuracy: 0.7511209166607359
Recall: 0.8037423225253535
Precision 0.7260645161290322
ROC - AUC 0.751303785376152
F1 Score 0.762931326689716
```

```
In [68]: 1 # Add these values to our model dictionary
2 # Since pandas does not allow you to add rows without removing the indices correspond to the models,
3 # we need to recreate the table again
4
5 model_name = 'Neural Network'
6 model_arr[model_name] = nn_model
7
8 train_acc[model_name] = nn_trn_acc
9 test_acc[model_name] = nn_tst_acc
10 rec[model_name] = nn_rec
11 prec[model_name] = nn_rec_prec
12 F1[model_name] = nn_rec_F1
13 Roc_Auc[model_name] = nn_rec_roc_auc
14 run_time[model_name] = 48 # Pulled from neural network notebook. This number is from the early stoppi
```

```
In [69]: 1 measures = pd.DataFrame(index=model_arr.keys(), columns=['Training Accuracy', 'Testing Accuracy', 'Recall', 'Precision', 'F1 Score', 'Roc-AUC Score', 'Runtime (s)'])
2 measures['Training Accuracy'] = train_acc.values()
3 measures['Testing Accuracy'] = test_acc.values()
4 measures['Recall'] = rec.values()
5 measures['Precision'] = prec.values()
6 measures['F1 Score'] = F1.values()
7 measures['Roc-AUC Score'] = Roc_Auc.values()
8 measures['Runtime (s)'] = run_time.values()
9 measures
```


Out[69]:

	Training Accuracy	Testing Accuracy	Recall	Precision	F1 Score	Roc-AUC Score	Runtime (s)
Logistical Regression	0.745076	0.749413	0.769187	0.738474	0.753518	0.749493	0
Random Forest	0.997313	0.742652	0.780620	0.724115	0.751307	0.742805	7
Decision Tree Classifier	0.997331	0.655327	0.653566	0.654033	0.653799	0.655320	0
XGB Classifier	0.791214	0.748345	0.788767	0.728389	0.757376	0.748509	5
SVC	0.769719	0.753541	0.806917	0.727765	0.765300	0.753756	350
GaussianNB	0.714614	0.717743	0.703016	0.722638	0.712692	0.717683	0
KNeighbors	0.796107	0.709487	0.734886	0.697788	0.715857	0.709589	275
XGB Tuned 1	0.757567	0.754466	0.798199	0.732651	0.764022	0.754643	5
XGB Tuned 2	0.764435	0.753897	0.797485	0.732187	0.763442	0.754073	38
Neural Network	0.757800	0.759875	0.788481	0.744434	0.765825	0.759990	48

```
In [70]: 1 TN_nn, FP_nn, FN_nn, TP_nn = confusion_matrix(y_test, y_pred_nn_rnd).ravel()
2
3 print('True Positive(TP) = ', TP_nn)
4 print('False Positive(FP) = ', FP_nn)
5 print('True Negative(TN) = ', TN_nn)
6 print('False Negative(FN) = ', FN_nn)
```

```
True Positive(TP) = 5517
False Positive(FP) = 1894
True Negative(TN) = 5160
False Negative(FN) = 1480
```

```
In [71]: 1 cm_nn = confusion_matrix(y_test, y_pred_nn_rnd)
2
3 cm_nn = ConfusionMatrixDisplay(confusion_matrix = cm_nn, display_labels = [False, True])
4
5 cm_nn.plot()
6 plt.show()
```

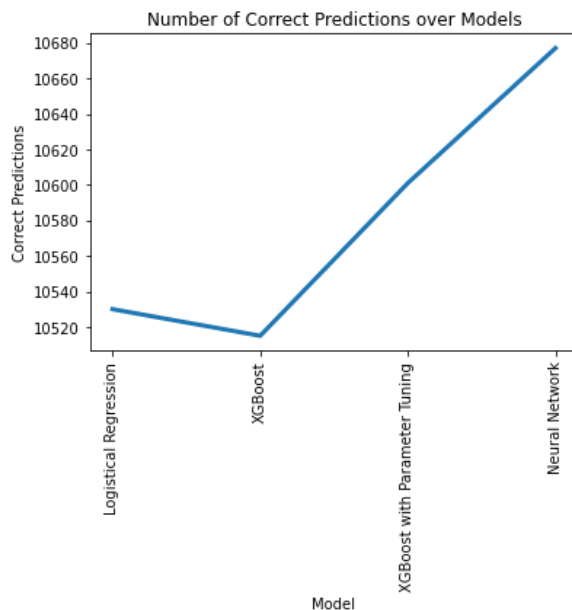


The neural network identified slightly more true positives and true negatives.

```
In [72]: 1 # Find the difference in correct predictions between all models
2 # Correct Predictions are defined as the number of TP + TN
3
4 lr_corr_pred = lr_TP + lr_TN # Correct number of predictions made by baseline logistic regression model
5 xgb_corr_pred = xgb_TP + xgb_TN # Correct number of predictions made by XGBoost model
6 xgb_tnd_corr_pred = TP_xgb_tnd + TN_xgb_tnd # Correct number of predictions made by XGBoost tuned model
7 nn_corr_pred = TP_nn + TN_nn # Correct number of predictions made by neural network
8
9 diff_preds_1 = xgb_corr_pred - lr_corr_pred
10 diff_preds_2 = xgb_tnd_corr_pred - xgb_corr_pred
11 diff_preds_3 = xgb_tnd_corr_pred - lr_corr_pred
12 diff_preds_4 = nn_corr_pred - xgb_tnd_corr_pred
13
14
15 print("The initial XGBoost model made",diff_preds_1,"more correct predictions than the baseline model")
16 print("The tuned XGBoost model made",diff_preds_2,"more correct predictions than the initial XGBoost model")
17 print("The tuned XGBoost model made",diff_preds_3,"more correct predictions than the baseline model.")
18 print("The neural network made",diff_preds_4,"more correct predictions than the tuned XGBoost model.")
```

The initial XGBoost model made -15 more correct predictions than the baseline model.
 The tuned XGBoost model made 86 more correct predictions than the initial XGBoost model.
 The tuned XGBoost model made 71 more correct predictions than the baseline model.
 The neural network made 76 more correct predictions than the tuned XGBoost model.

```
In [73]: 1 x_axis = ["Logistical Regression", "XGBoost", "XGBoost with Parameter Tuning", "Neural Network"]
2 y_axis = [lr_corr_pred, xgb_corr_pred, xgb_tnd_corr_pred, nn_corr_pred]
3
4 plt.plot(x_axis, y_axis, linewidth = 3)
5 plt.xlabel('Model')
6 plt.ylabel('Correct Predictions')
7 plt.title('Number of Correct Predictions over Models')
8 plt.xticks(rotation=90)
9
10 # Show the plot
11 plt.show()
```



Interestingly, like the paper the neural network is the most accurate model. We did not seem the same level of difference in accuracy that was observed in the paper (82% vs 79%), however, we may not have the computing resources to add more layers and create a denser neural network.

```
In [74]: 1 # Calculate percentage increase in accuracy between the most accurate model and the least accurate
2
3 print("The neural network is", round((nn_corr_pred - lr_corr_pred)/diab_df.shape[0] * 100, 2), "% more accurate than the base model")
```

The neural network is 0.21 % more accurate than the base model

Though iterative modeling, we've improved the efficacy of our model by 125 predictions. This represents an increase of around ~0.2%.

Final Model Evaluation

For the final model, we are recommending our baseline model the logistic regression model for use by the CDC. Even though we iteratively improved the accuracy and precision metrics across the XGB, tuned models, and neural network, the increase in these metrics is not worth the time and resources it takes to train and tune these models.

We only used a small sample of the data available in the study in our training/testing (31K vs 440K records). Deploying this model across data that goes into the hundreds of thousands or millions of records if you consider the previous years data, may not be economical if you consider time, computational resources, and FTEs it takes.

Logistical Regression probably gives the CDC what they need to reasonably determine the likelihood of diabetes with a limited budget.

In [75]: 1 measures

Out [75]:

	Training Accuracy	Testing Accuracy	Recall	Precision	F1 Score	Roc-AUC Score	Runtime (s)
Logistical Regression	0.745076	0.749413	0.769187	0.738474	0.753518	0.749493	0
Random Forest	0.997313	0.742652	0.780620	0.724115	0.751307	0.742805	7
Decision Tree Classifier	0.997331	0.655327	0.653566	0.654033	0.653799	0.655320	0
XGB Classifier	0.791214	0.748345	0.788767	0.728389	0.757376	0.748509	5
SVC	0.769719	0.753541	0.806917	0.727765	0.765300	0.753756	350
GaussianNB	0.714614	0.717743	0.703016	0.722638	0.712692	0.717683	0
KNeighbors	0.796107	0.709487	0.734886	0.697788	0.715857	0.709589	275
XGB Tuned 1	0.757567	0.754466	0.798199	0.732651	0.764022	0.754643	5
XGB Tuned 2	0.764435	0.753897	0.797485	0.732187	0.763442	0.754073	38
Neural Network	0.757800	0.759875	0.788481	0.744434	0.765825	0.759990	48

Recommendations

- The CDC should use the logistical regression model in their application.
- Consider a strategy around educating people to take their blood pressure on a regular basis since it was one of the top features.
- Providers who see people with high cholesterol should also screen for diabetes since high cholesterol was another top feature.
- Continue advocating for policy/strategies that aim to improve the general health and fitness of Americans. Low health was the most correlated feature with diabetes.

Future Projects

- Evaluate previous BRFSS data sets. Measure the rate of diabetes and other chronic conditions to find their trends across the country.
- Use the model to create an application on the CDC's website that allows a person to enter their data and get a diabetic risk score.
- Further investigate a strategy around making it easier for people to take and track their blood pressure. It was found to be the greatest predictor around diabetes.

Reproduction Steps

Download from Github to Local Machine

1. Download the 2015.CSV from this link: <https://www.kaggle.com/datasets/cdc/behavioral-risk-factor-surveillance-system> (<https://www.kaggle.com/datasets/cdc/behavioral-risk-factor-surveillance-system>)
2. Save CSV to file and run steps from the data cleaning notebook
3. Run the main notebook.